

Overview of Microprocessor Designs



Chun-Jen Tsai
NYCU
10/6/2023


General Purpose Computers

- ❑ The first electronic general-purpose (programmable) computer is ENIAC
 - Built in 1946
 - Programmed by re-wiring functional boxes

- ❑ The first instruction-based computer is EDVAC
 - Realization of the “stored-program concept”
 - Proposed by Eckert et al.
 - Program bits stored in memory, just like data
 - EDVAC built in 1948 by configuring an ENIAC
 - von Neuman wrote the first ISA draft for EDVAC

UNIVAC I

- ❑ UNIVAC I is the first commercial ISA-based computer
 - Build by Eckert-Mauchly Computer Corporation in 1951
 - Performance: 0.0008 DMIPS/Mhz



† Picture taken by U.S. Census Bureau employees - <https://www.census.gov/history/>, Public Domain,
<https://commons.wikimedia.org/w/index.php?curid=61118834>

From Architecture to Realization

- ❑ Architecture (a.k.a. ISA)
 - The design of instructions, addressing modes, and system interface of a microprocessor
- ❑ Implementation (a.k.a. microarchitecture)
 - The design of the circuit modules that implements the ISA
 - The high-level design description is mainly based on diagrams (hierarchical, timing, and interface diagrams, ..., etc.)
 - For each bottom circuit module, the description are often synthesizable RTL code (Verilog, VHDL)
- ❑ Realization
 - Mapping of the synthesizable circuit description to a specific target technology (e.g., FPGAs or ASICs)

CISC vs RISC ISA

- ❑ In computer history, we have Complex Instruction Set Computer (CISC) before Reduced Instruction Set Computer (RISC)
 - What defines “complex” operation?
 - Why did we build “complex” computers before “simple” computers?
 - Which one is “better” (consider Intel i9 vs ARM Cortex A78)?
 - Does instruction set architecture design really matter?

Scalar ISA vs Vector ISA

- A vector ISA packs several numbers in a long bit-width register and perform the same operation on all data items concurrently, for example,

```
rv1 ← [a, b, c, d]; // 4 32-bit int, rv1 is of 128 bits  
rv2 ← [e, f, g, h];  
rv3 ← rv1 + rv2;    // [a+e, b+f, c+g, d+h]
```

- Intel SSE, and the infamous AVX-512 instructions are vector extensions to the base scalar ISA
- ARM defines scalar instructions in the base ISA, and vector instructions in the NEON coprocessor ISA

ISA Design

- ❑ ISA is not that important for modern computers
 - ISA designs have converged over several decades
 - Basic ISAs for different processors show little differences
 - Application-specific instructions are the main differentiators

- ❑ ISA does impact the microarchitecture complexity
 - The Dynamic/Static Interface[†] (DSI) determines what features of the processor are implemented in HW or SW
 - An operation can be implemented by compiler/programmer using instructions (dynamic) or wired into the processor (static)
 - Software interpreters and hardware microprogramming blurs the boundary of DSI

[†] S. W. Melvin and Y. N. Patt, "A clarification of the dynamic/static interface," Proc. of the 12th Hawaii Int. Conf. on System Science, 1987.

DSI of a Processor

- ❑ DSI provides an important separation between the architecture and the implementation
 - A CISC ISA places the DSI at the higher-level assembly language, and hardwired complex operations
 - A RISC ISA lowers the DSI and expects to perform more of the optimizations above the DSI via the compiler
- ❑ DSI decides the microarchitecture features of a μP
 - Java bytecode has an instruction for multi-dimensional memory allocation!
 - Today, the DSI is chosen towards a simpler ISA architecture

Smart ISA Designs that Impact HW

❑ Instructions with additional barrel shifter


- In ARM, two ALU ops in 1 cycle:

```
ADD r3, r2, r1, LSL#3 ; r3 ← r2 + r1 * 8
```

❑ Conditional executions

- In ARM, all ALU ops are conditional:

```
CMP r0, #5
BEQ Bypass ; if (r0 != 5)
ADD r1, r1, r0 ; r1 ← r1 + r0
SUB r1, r1, r2
Bypass . . .
```



→

```
CMP r0, #5
ADDNE r1, r1, r0
SUBNE r1, r1, r2
```

❑ Default branch direction bit

- Compilers can encode the default jump direction in a loop
- Example: PowerPC ISA

Smart ISA Designs Become Obsolete

- ❑ With new hardware design practices, these smart ISA designs may not be useful:
 - Barrel shifters along operand path hinder clock rates
 - Conditional execution still wastes CPU cycles
 - Default branch bits do not account for dynamic behavior
- ❑ In particular, these tricks are meaningless under superscalar microarchitecture
- ❑ For ISA design, simplicity is the king!

Design Objectives of Processors

- ❑ There are several design objectives of a μ P:
 - Performance
 - Power
 - Cost
 - Reliability
 - Security

- ❑ Which one has more impact on achieving these goals:
ISA or microarchitecture?

Pursuit of Performance

- The factors that affect the processor performance:

$$\text{Execution time} = \frac{\text{instruction}}{\text{program}} \times \frac{\text{cycle}}{\text{instruction}} \times \frac{\text{time}}{\text{cycle}}$$


Compiler & OS

CPI: μ Arch


μ Arch & fabrication tech.

Stack vs. Register Architecture

- ❑ A processor can adopt a stack- or register-based microarchitecture depending on where the intermediate results are stored:




Stack-based datapath



Register-based datapath

VLIW Architecture

- ❑ Very Long Instruction Word (VLIW) ISA allows multiple operations to be packed into a bundle for execution:
 - Popular for DSP processors
 - Each bundle has a fixed number of “issue slots”
 - All operations in a bundle issued at the same clock cycle
 - Each issue slot encodes one instruction

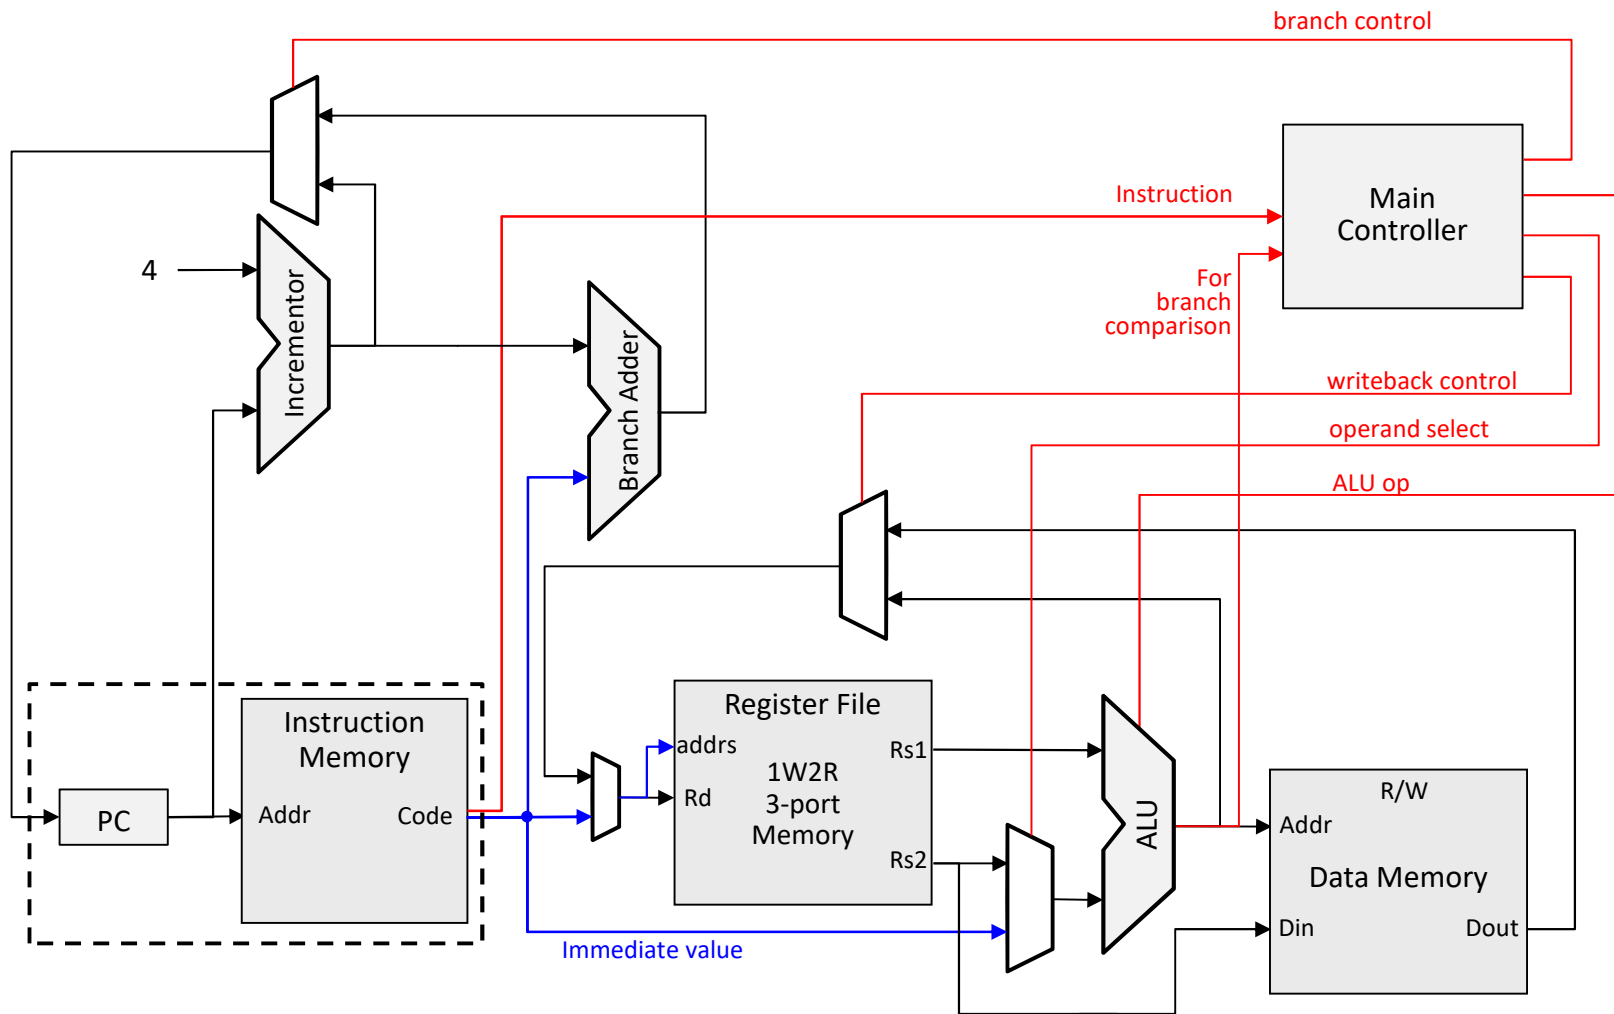


- Easier to program by using “C intrinsics”

Scalar vs Superscalar Architecture

- ❑ Traditional scalar processors “execute” one instruction per clock cycle
- ❑ Superscalars execute multiple instructions per cycle:
 - Need multiple ALU functional units
 - Need multiple read/write buffers (registers)
 - Need internal queues for decoded instructions
 - Need control schemes to enforce data-dependency
 - Maybe in-order or out-of-order for instruction execution

1-Cycle Simple Scalar μ Architecture



Multiple Issue of Instructions

❑ Static multiple issue

- Typically for VLIW architecture
 - Compiler/programmer groups instructions to be issued together
 - Packages them into “issue slots”
 - Compiler/programmer detects and avoids hazards
- VLIW compilers are very difficult to design

❑ Dynamic multiple issue

- CPU chooses instructions to issue each cycle
 - Execution can be in-order or out-of-order
- Compiler can also help by reordering instructions
- CPU resolves hazards using advanced techniques at runtime

Speculative Execution of Instructions

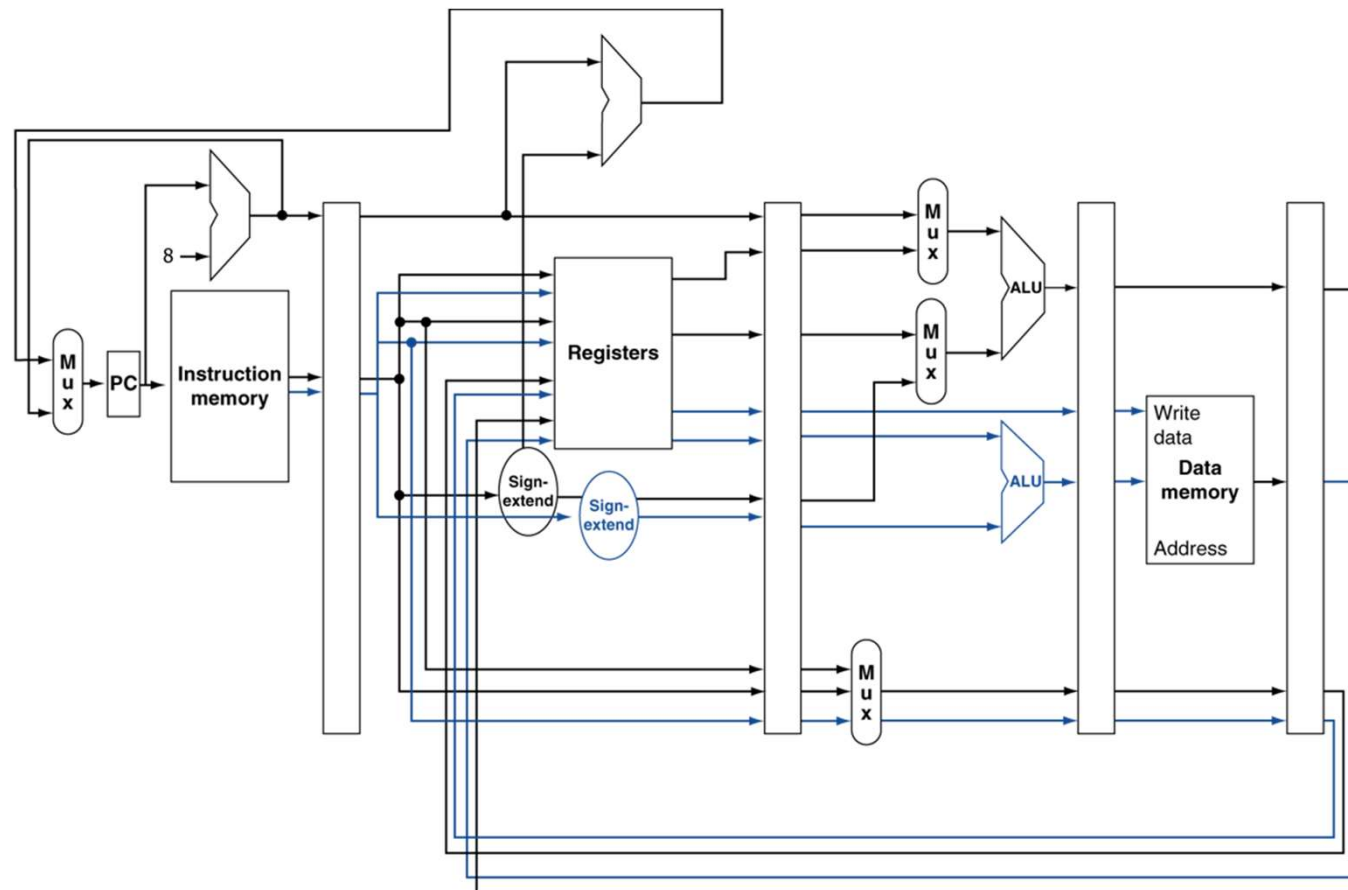
- ❑ Speculate on branch
 - Execute a branch instruction before the outcome is known
 - Roll back if path taken is different

- ❑ Speculate on data loading
 - Execute an indirect load before the address is determined
 - Pointer in register may be wrong due to out-of-order execution
 - Roll back if location is updated

- ❑ In general, can look ahead for instructions to execute
 - Buffer results until they are actually needed (for writeback)
 - Flush buffers on incorrect speculation

Static Double-Issue μ Architecture

- ❑ Two instructions will be concurrently executed:
 - Only one of them can be load/store



Dynamic Multiple Issue

- ❑ Superscalar processors
 - Instruction “grouping” is done by the CPU control logic
 - Execution can be in-order or out-of-order
 - Avoiding structural and data hazards

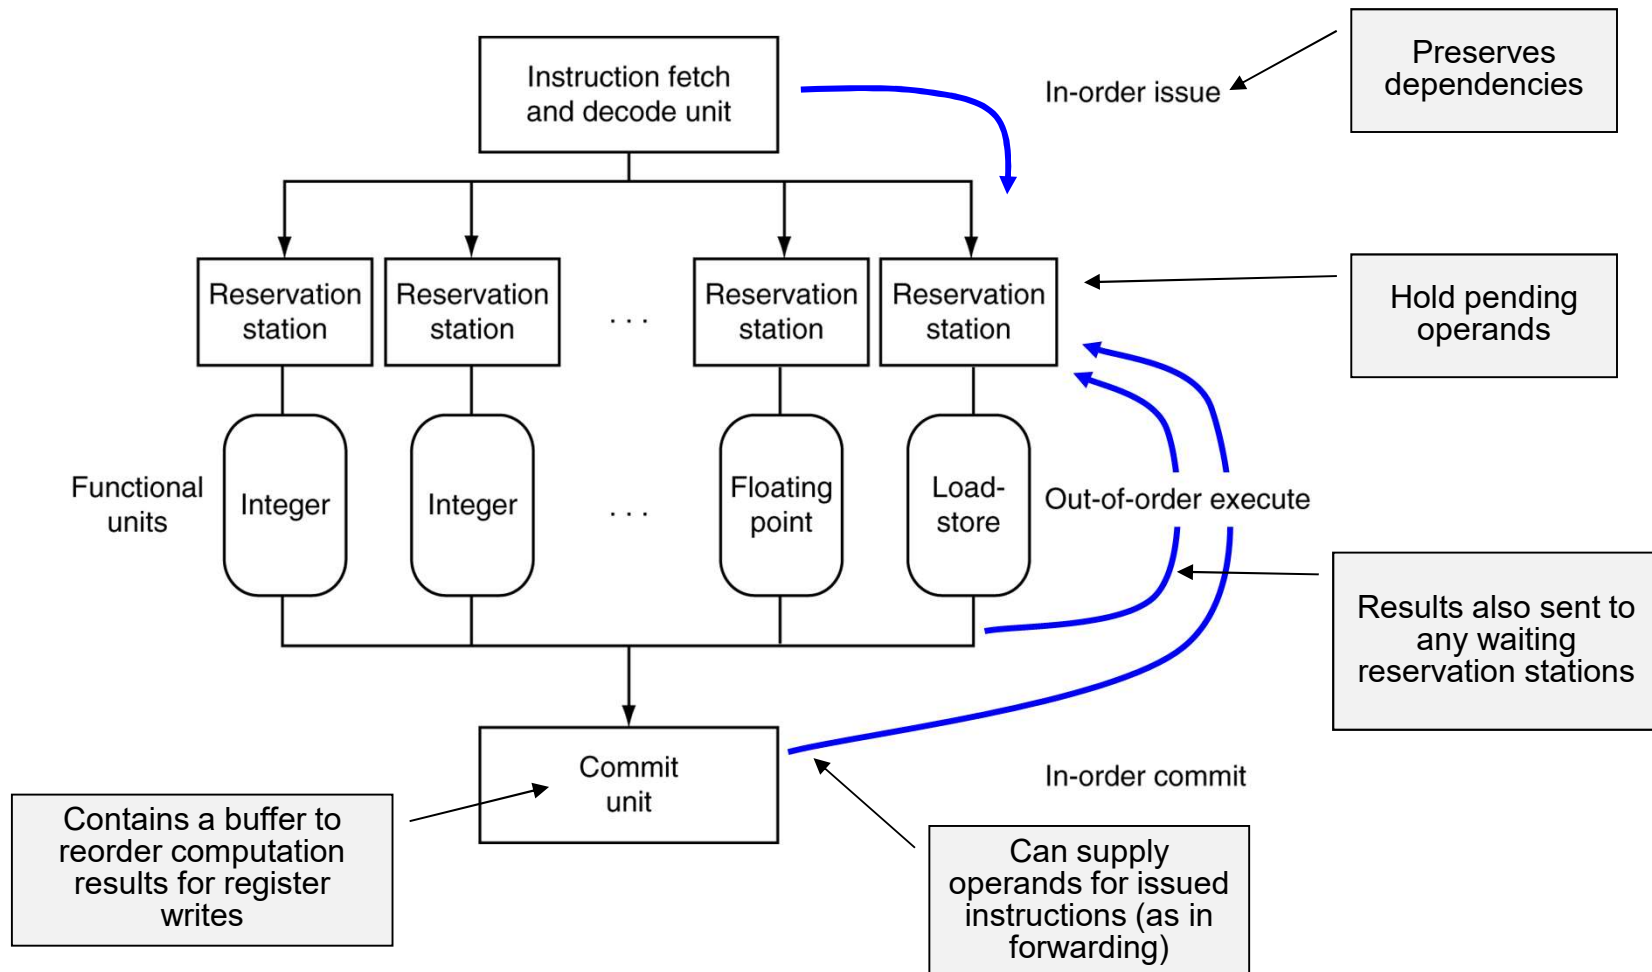
- ❑ Avoids the need for compiler scheduling
 - Compilers may still help
 - CPU analysis windows are small
 - Roll-backs are expensive
 - Code semantics ensured by the CPU
 - CPU has to ensure the results are precisely what the code says

Dynamic Pipeline Scheduling

- ❑ Allow the CPU to execute instructions out-of-order to avoid stalls
 - Issue instructions in order
 - Execute instructions out of order
 - Commit results to registers in order
- ❑ Example: CPU can start `sub` and `slti` while `add` is waiting for `lw`

```
lw    $t0, 20($s2)
add   $t1, $t0, $t2
sub   $s4, $s4, $t3
slti  $t5, $s4, 20
```

Dynamically Scheduled CPU



Register Renaming

- ❑ Typical ISA has a limited amount of registers (e.g. 32)
 - Architectural register number is limited by instruction coding
 - CPU can have more physical registers
 - An architectural register can be renamed to a physical register

- ❑ Reservation stations and reorder buffer effectively provide register renaming upon instruction issue
 - If operand is available in register file or reorder buffer
 - Copied to a reservation station
 - No longer required in the register; can be overwritten
 - If operand is not yet available
 - It will be provided to the reservation station by a functional unit
 - Register update may not be required

Speculation in Reservation Station

❑ Branch speculation

- Predict branch and continue issuing
- Don't commit results until branch outcome determined

❑ Load speculation

- Avoid load and cache miss delay
 - Predict the effective address
 - Predict loaded value
 - Load before completing outstanding stores
 - Bypass stored values to load unit
- Don't commit load until speculation cleared

Name-Dependence Hazards

- ❑ Name hazards of out-of-order, multiple-issue execution:
 - Two instructions use the same register name/memory location
- ❑ Data and name dependencies are classified as follows:
 - Write-After-Read (WAR) hazard: an anti-dependence of name

```
add    $t1, $t2, $t3
add    $t3, $t0, $t1
```

- Write-After-Write (WAW) hazard: an output name dependence

```
add    $t3, $t2, $t1
add    $t3, $t0, $t1
```

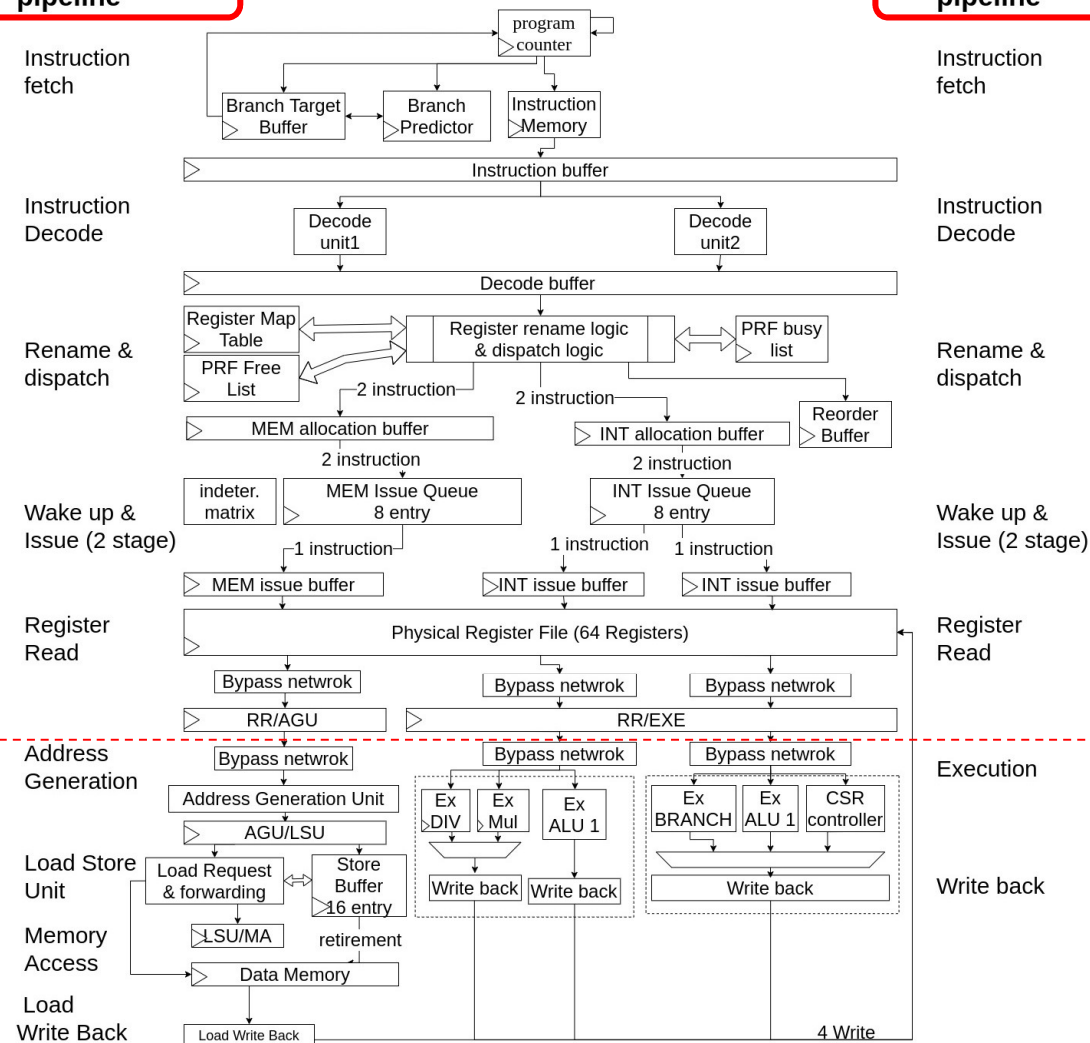
- Read-After-Write (RAW) hazard: a true data dependence

```
add    $t3, $t0, $t1
add    $t1, $t2, $t3
```

Superscalar Example (Falco)

Memory Access pipeline

Computation pipeline



Evolution History of ARM μ Arch.

- ❑ Acorn Computers, established in 1978
 - ARM1 (1985) → ... → ARM3 (1989)

- ❑ ARM, established in 1990
 - Classical ARMs:
 - ARM6 (1992) → ARM7 (1993) → ARM9 (1998) → ARM11 (2002)
 - ARM6 was used in the Apple Newton PDA (1993)

 - Cortex ARMs:
 - Cortex A8(2005) → ... → Cortex A78(2020) → Cortex A710(2021)
 - Cortex M3(2004) → ... → Cortex M55(2020)
 - Cortex X1(2020) → Cortex X2(2021)

Pipeline Stage Design Choices

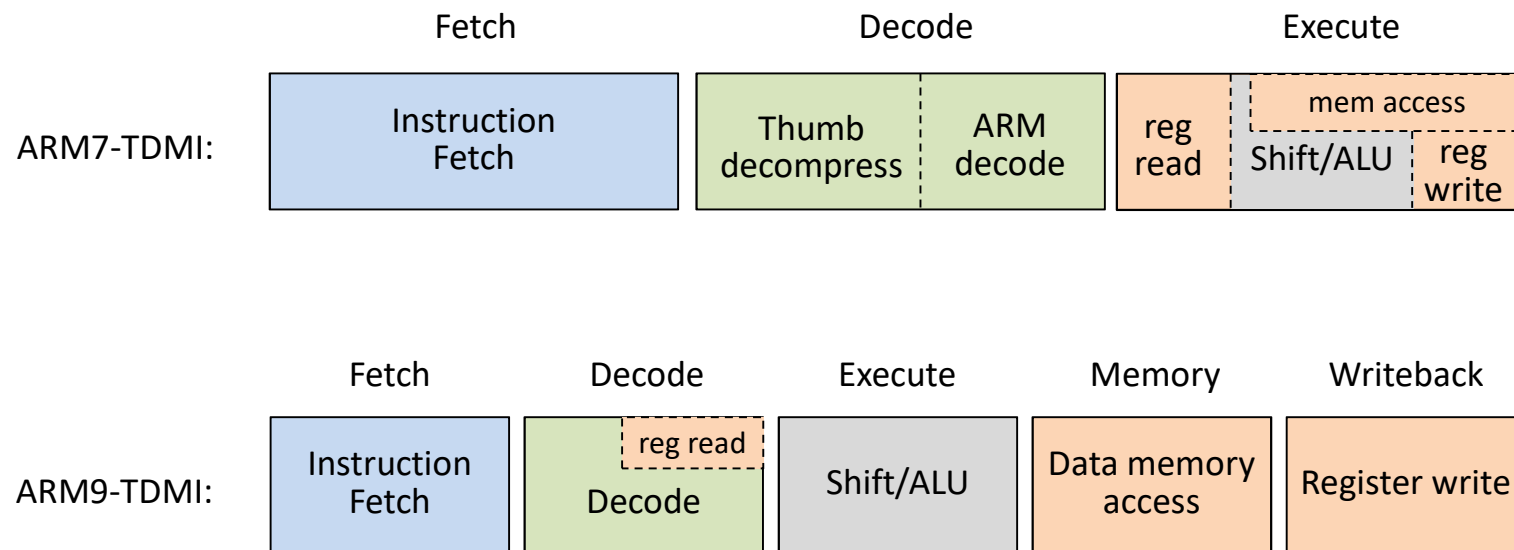
- ❑ In previous design, the “processor state” are stored in the PC and the register file
 - A 2-stage pipeline seems a natural implementation (assuming the instruction/data memory being latch-like combinational memory devices)
 - A synchronous memory would be a “state” element too
- ❑ Back in 1980’s, it is customary to cut the processor pipeline into three stages
 - Example: ARM 7 (ARMv4 architecture)
- ❑ Modern in-order execution RISC processors contains at least five pipeline stages

Classical RISC Pipeline Stages

- ❑ The execution of an instruction has several steps
 - Not all instructions require the same number of steps – a big problem for CISC processors
- ❑ For a RISC processor the classical steps are:
 - Fetch – fetch instruction from instruction memory
 - Decode – decode instruction to generate control signal
 - Execute – execute ALU operations
 - Memory – access data memory
 - Writeback – write results back to register file

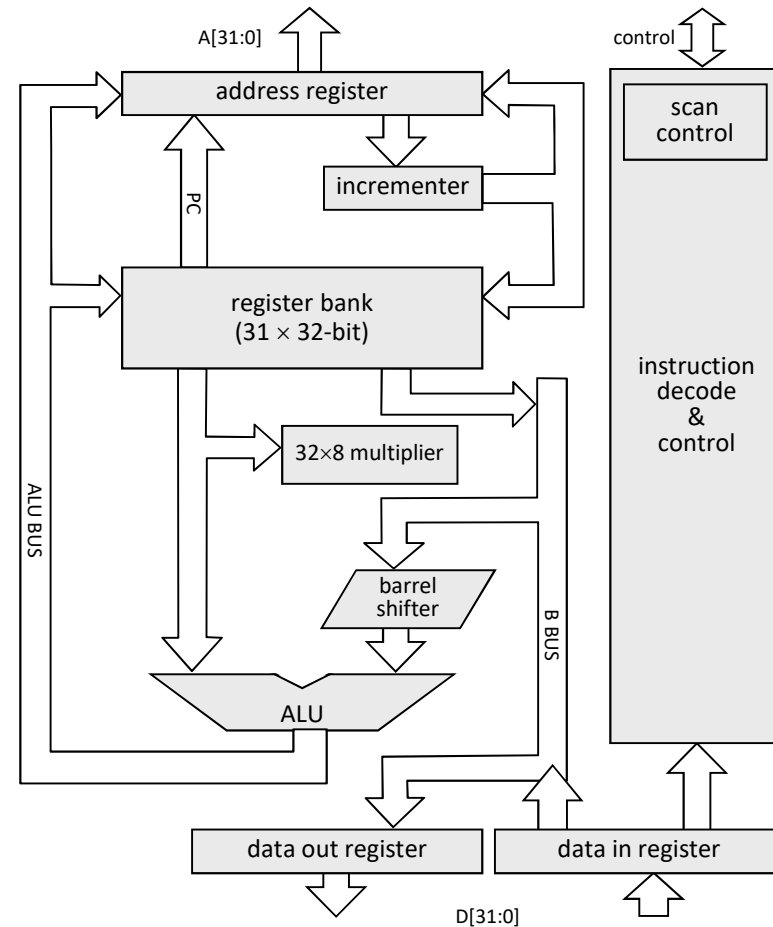
ARM Cores: 3-Stage vs. 5-Stage

- ❑ Pipeline stage organizations of a 3-stage ARM7 versus a 5-stage ARM9:




Classical ARM 7 μ Architecture

- ❑ One of the most famous three-stage pipeline RISC
 - ARM ISA v4T
 - Von Neumann architecture
 - ~ 80% market share in 2G cell phones
 - The predecessor of ARM Cortex-M μ Architecture
 - Used in the original iPod (customized edition)



- ❑ First 5-stage pipeline from ARM
- ❑ Harvard architecture
- ❑ Add facilities to remove pipeline hazard

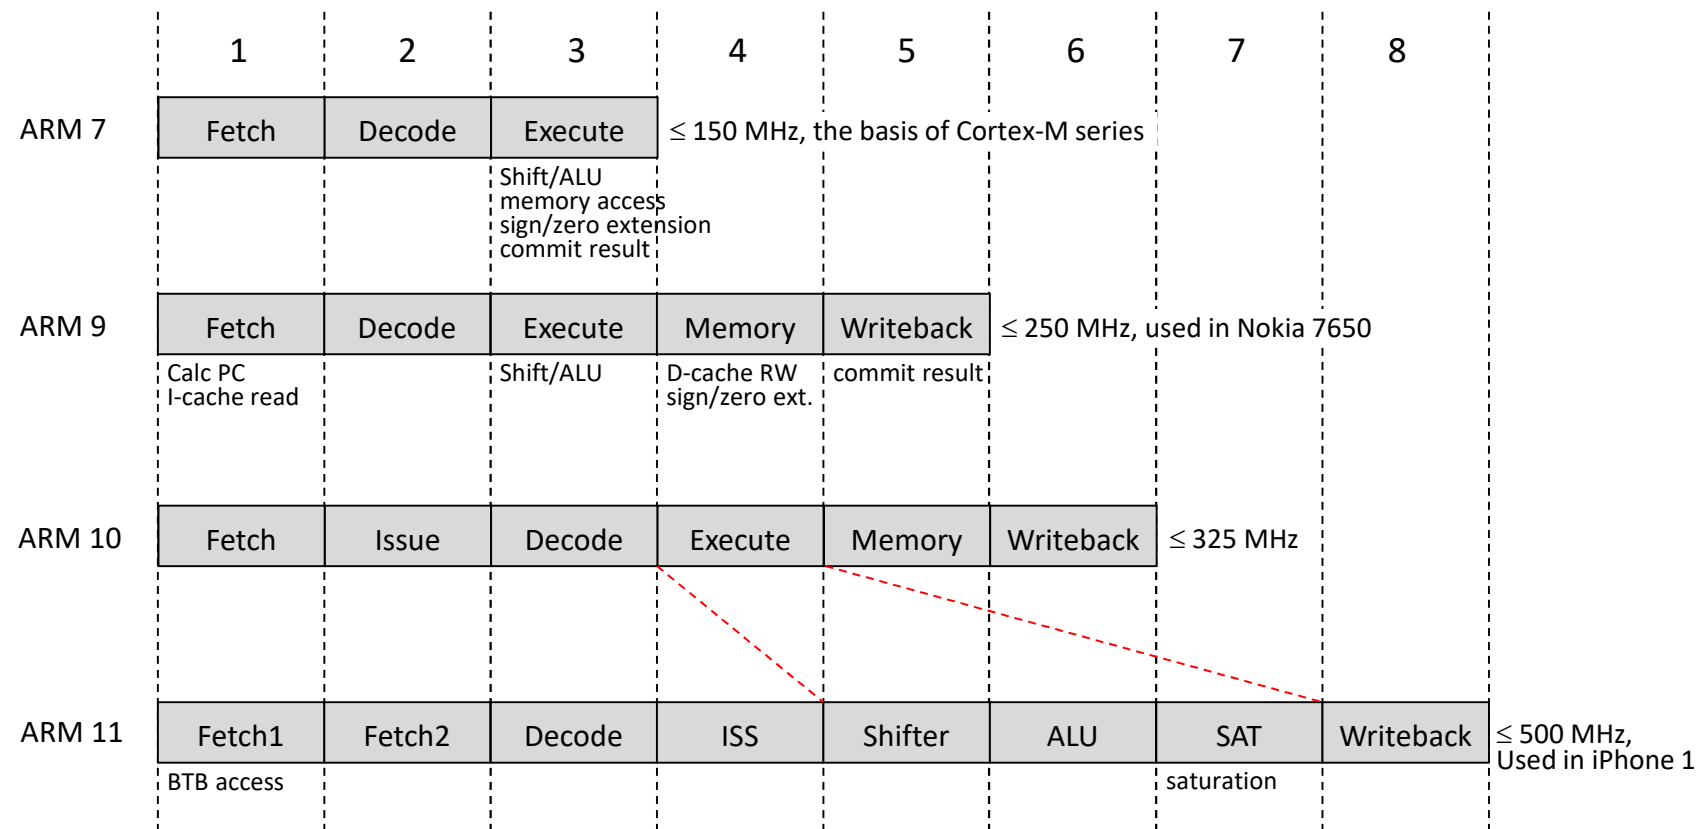


ARM 11 μ Architecture

- ❑ The final version of the classical ARMs
- ❑ Key μ Architecture features:
 - In-order execution, out-of-order completion (certain ops)
 - Non-blocking cache
 - Parallel ALUs
 - Parallel Loads/Stores

Classical ARM Pipeline Evolutions

❑ ARM classical pipelines before Cortex-A series:



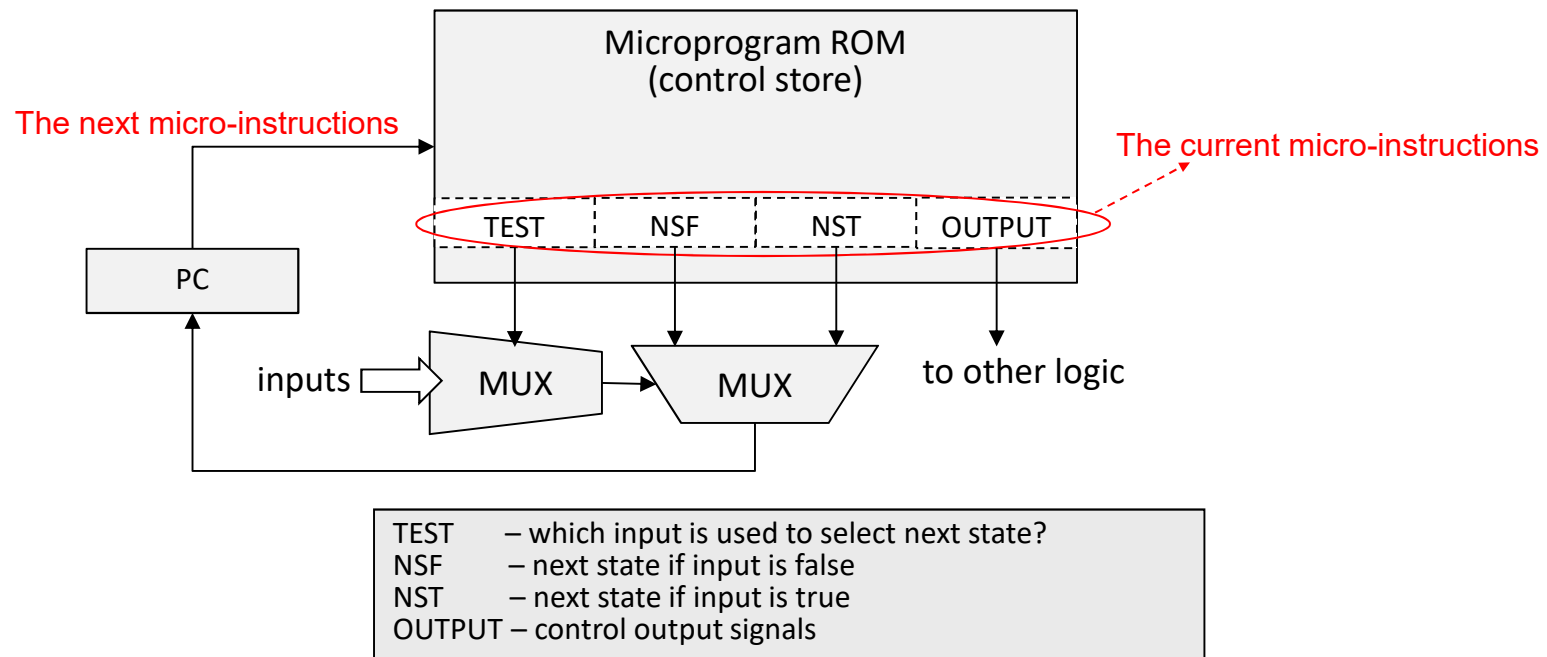
Microprogramming for CISC

- ❑ Hardwiring the controller FSM to execute a multi-cycle CISC ISA instruction is expensive
 - Design/debugging cost is quite high

- ❑ A more flexible approach is to use microprogramming[†]
 - Using a tiny (in-circuit) computer to execute the FSM
 - Each microinstruction specifies
 - The outputs to be generated
 - Where to find the next microinstruction (i.e. state transitions)
 - Moore FSMs are more suitable for microprogramming

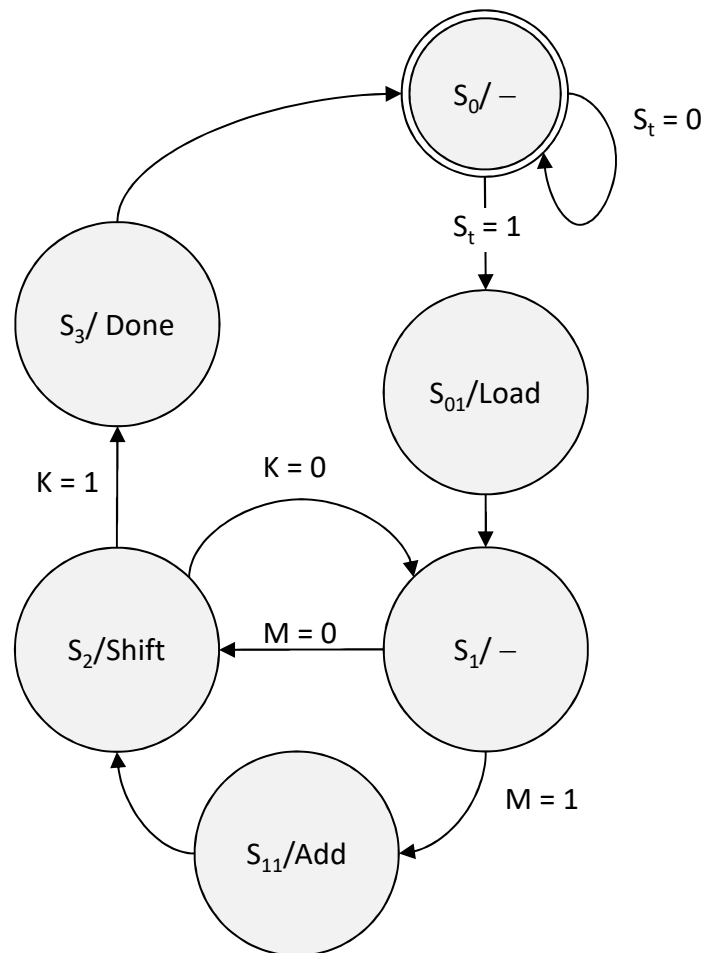
Microcode Sequencer

- ❑ Microprogrammed controllers are called sequencers
 - The control signals and the next state information are stored in a memory (i.e. ROM) inside the controller
 - The memory is called “control store” or “microcode memory”



Example: Multiplier Controller

- ❑ The FSM and its microprogram of a 4-bit multiplier




- Microprogram:

State	ABC	TEST	NSF	NST	Load	Add	Shift	Done
S_0	000	00	000	001	0	0	0	0
S_{01}	001	11	010	010	1	0	0	0
S_1	010	01	100	011	0	0	0	0
S_{11}	011	11	100	100	0	1	0	0
S_2	100	10	010	101	0	0	1	0
S_3	101	11	000	000	0	0	0	1

output signals

- Input select multiplexer:



Discussions

- ❑ Most smartphone processors are based on the ARM ISAs, but the microarchitecture may be different
 - Apple A/M-series: fully developed by Apple Inc.
 - Qualcomm Kryo: some are developed by Qualcomm
 - Samsung Exynos: some are developed by Samsung
 - ARM Cortex A/X-series: developed by ARM, used by MediaTek, Huawei, Qualcomm, Samsung, and many other companies
→ where are the differences ?

- ❑ With the arrival of the RISC-V ISA, companies will design their own CPUs for marketing differentiation