

Application Processors and the Aquila SoC



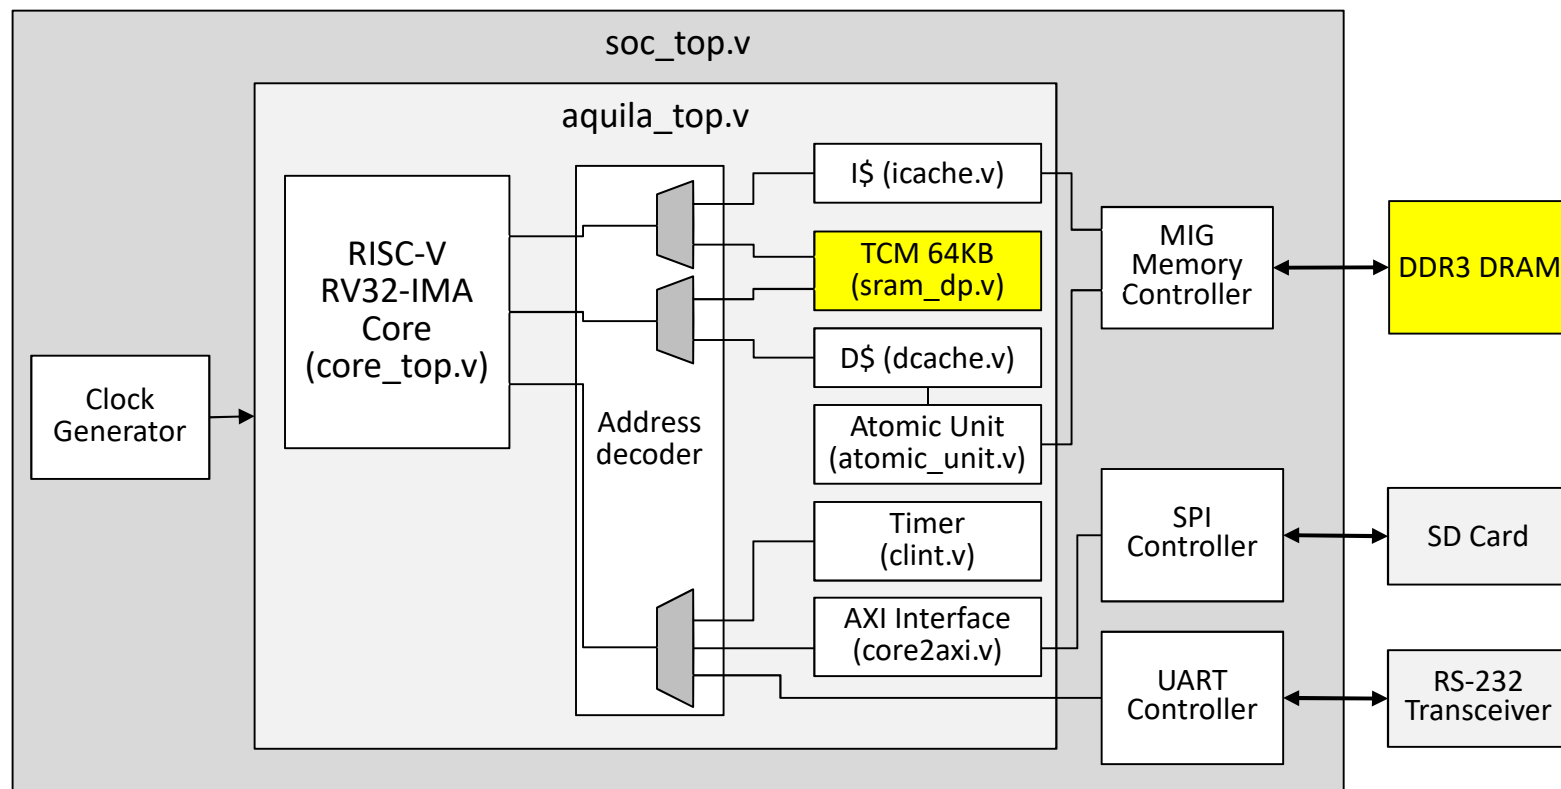
Chun-Jen Tsai
NYCU
10/27/2023

Aquila – An Open Source RISC-V

- ❑ The spec of the complete version of Aquila SoC:
 - RV32IMA + Zicsr + Zifencei ISA-compliant
 - Embedded 64KB tightly-coupled on-chip memory (TCM)
 - L1 4-way set associative data and instruction caches
 - CLINT for standard timer interrupts
 - CSRs and complete M-, U-, and S-mode support
 - SV32 Memory Management Unit
 - SD card I/O support
 - Ethernet support
 - The RTL model written in Verilog


SoC Interfaces (1/2)

- ❑ Aquila core can be used as:
 - A reusable IP (using AXI4 bus protocol)
 - A proprietary SoC (using proprietary bus protocol)



SoC Circuit Block Interfaces (2/2)

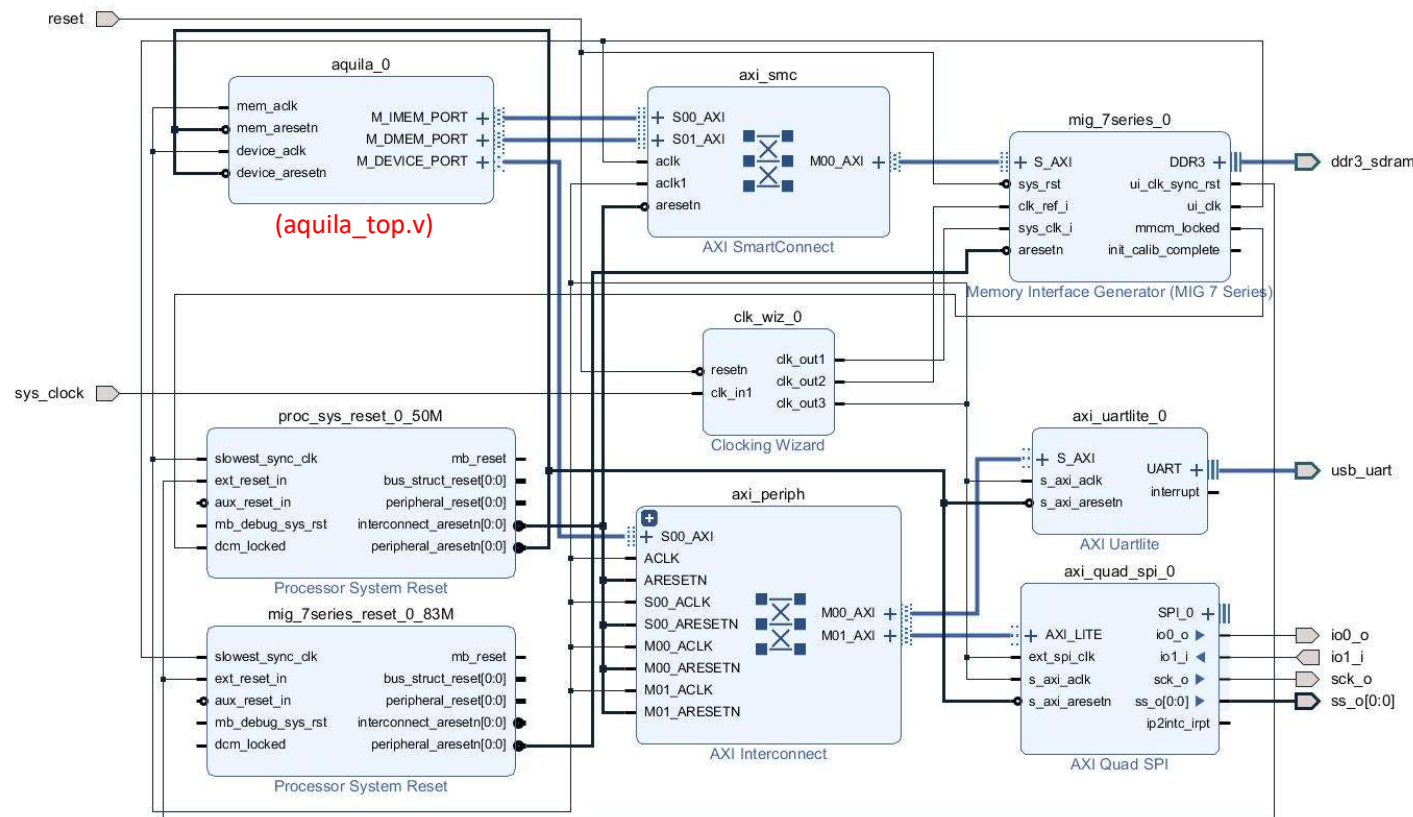
- ❑ Aquila SoC uses several interfaces to integrate slave circuit modules:



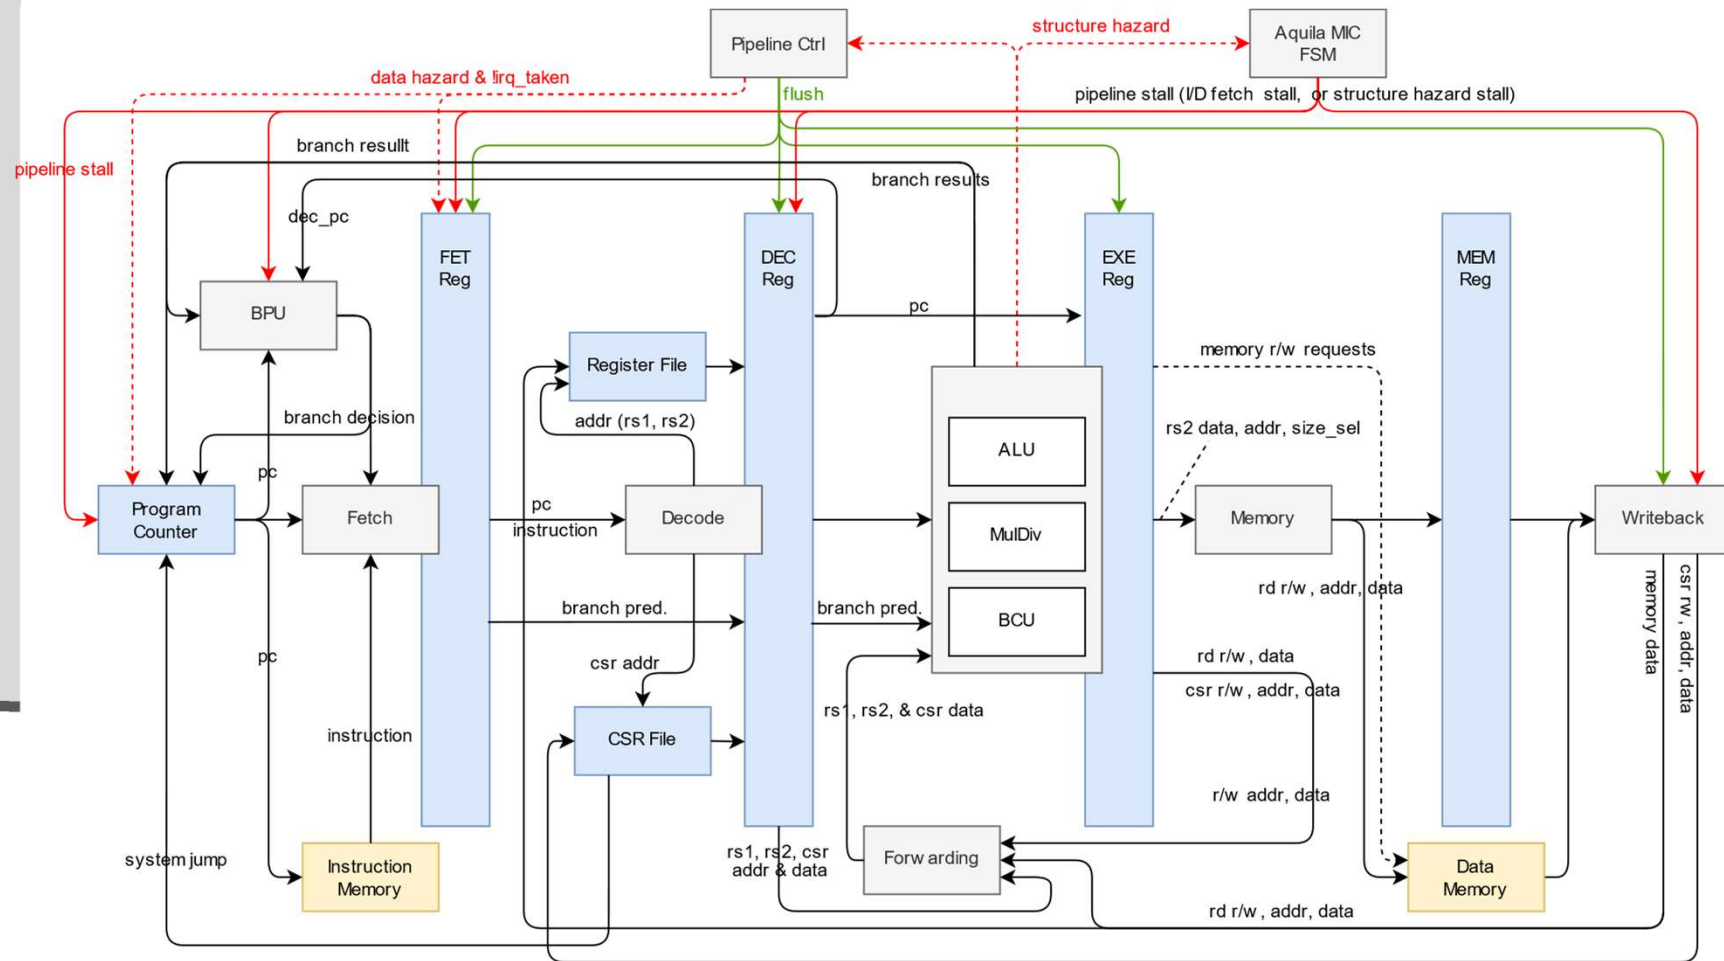
- ❑ Custom logic interfaces
 - Native MIG memory controller interface
 - cache controller interface
- ❑ General-purpose bus interfaces
 - Minimal synchronous handshaking bus: Timer, UART controller
 - Industry-standard bus (AXI): SPI controller for SD card

Aquila as an AXI Reusable IP

- ❑ Aquila can also be in a full-AXI infrastructure
 - AXI bus is designed by ARM to ease complex SoC integration
 - Vivado has an GUI-based IP Integrator tool for this purpose:




Aquila Pipeline Organization




Aquila Memory Interface Controllers

- ❑ Memory interface controllers (MIC) in `core_top.v` controls the instruction & data fetch states:



Instruction fetch controller

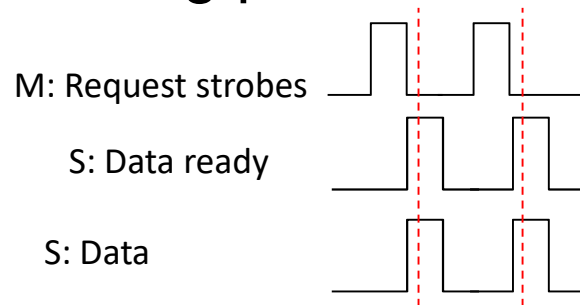
`stall = stall_instr_fetch || stall_from_exe`



Data access controller


Memory Access Signaling

- ❑ Aquila MIC uses the strobing protocol
- ❑ Strobing protocol: single-cycle request (read)



M means master, S means slave.

- ❑ Handshaking protocol: multi-cycle handshaking (read)



M means master, S means slave.

Pipeline Controller

- ❑ Pipeline Controller sends out flush signals and data hazard signals to different pipeline stages
- ❑ Different stages are flushed on different conditions:
 - Fetch – mis-branch, sys_jump[†], or fence_i
 - Decode – mis-branch, sys_jump, fence_i, load-use hazard, unsupported instruction
 - Execute – sys_jump or fence_i
 - Memory – No register to flush
 - Writeback – sys_jump

[†] sys_jump means an exception or interrupt happens!

Fetch Stage

- ❑ Receives instructions from the instruction memory, and pass it to the Decode unit

- ❑ Functions
 - Delay the instruction upon stall
 - Send “NOP” instruction upon flush, invalid instruction
 - Branch prediction is often considered as part of the Fetch stage tasks
 - Detect page fault exceptions when we have MMU
 - Pass exceptions to later stages

Decode Stage

- ❑ Extracts different fields from the instruction and sets control and immediate value signals accordingly
 - Very tedious but straightforward to implement

- ❑ Functions:
 - Decode instructions
 - Detect data hazard
 - Happens the 'rd' of the last instruction is the same as either the 'rs1' or 'rs2' of the current instruction
 - Pass exceptions to later stages

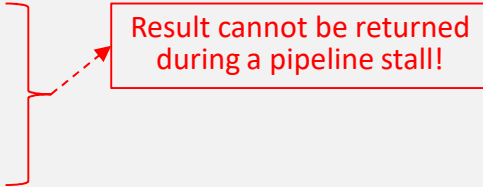
Execute Stage

- ❑ The Execute carries out data manipulation. The critical path of a processor usually involves this stage
 - On FPGAs, the critical path may get a little bit tricky
- ❑ Functions:
 - Single-cycle ALU operations
 - Multi-cycle ALU operations
 - Branch comparison calculations (using exclusive comparators)
 - Branch address computation (using exclusive adder)
 - Sent out memory read/write requests
 - May generate ALU exceptions
 - Pass exceptions to later stages

Multi-Cycle Instructions

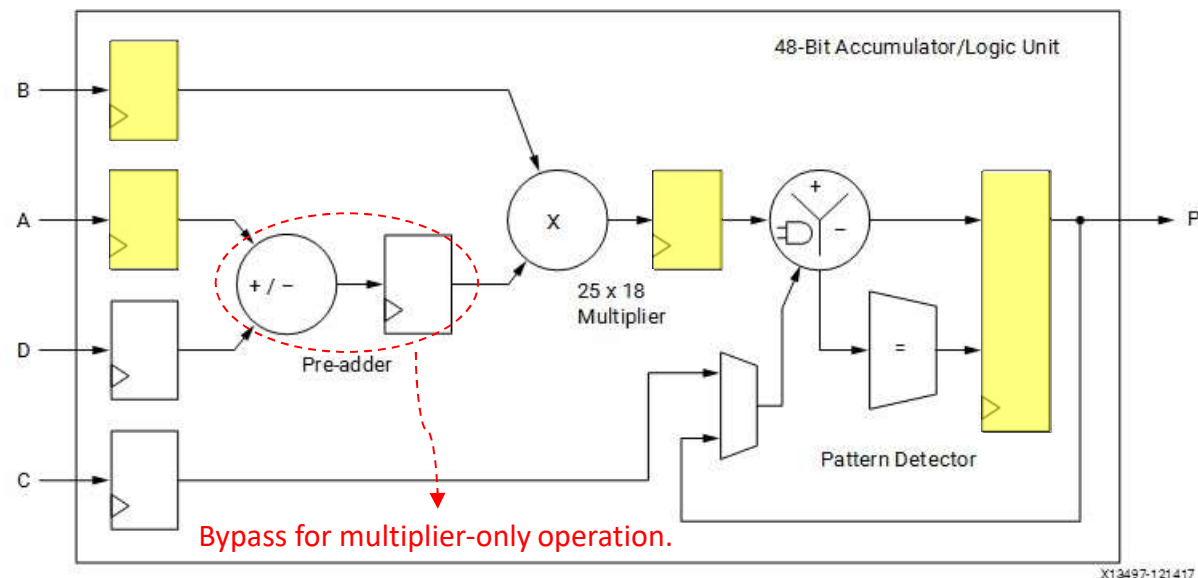
- ❑ RISC-V Ext. M is in the MulDiv module
 - This module is more of a template for multi-cycle instructions
 - The FSM for multi-cycle multiply/division:

```
always @(*)
begin
  case (S)
    S_IDLE:
      S_nxt = (req_i)? (is_a_zero | is_b_zero)? S_DONE : S_CALC : S_IDLE;
    S_CALC:
      S_nxt = (is_calc_done)? S_SIGN_ADJUST : S_CALC;
    S_SIGN_ADJUST:
      S_nxt = S_DONE;
    S_DONE:
      S_nxt = (stall_i)? S_STALL : S_IDLE;
    S_STALL:
      S_nxt = (stall_i)? S_STALL : S_IDLE;
    default:
      S_nxt = S_IDLE;
  endcase
end
```



Multiplier on FPGA (1/2)

- ❑ On Xilinx FPGAs, multiplier can be implemented using LUTs or DSP48 slices
 - Aquila has both implementations
- ❑ For DSP48 multiplier, you can specify the #stages
 - There are three sets of registers along the combinational path
 - Each set can be bypassed for a long combinational path



Multiplier on FPGA (2/2)

- ❑ To use DSP slices for 32×32 multiplication, four slices will be used:

$$((a \ll 16) + b) \times ((c \ll 16) + d) = (a \times c) \ll 32 + (a \times d) \ll 16 + (b \times c) \ll 16 + b \times d.$$

- ❑ The number of pipeline stages can be inferred using the following code patterns:

```
always @(posedge clk)
begin
    m_r <= v1*v2;
end
```

one-stage

```
always @(posedge clk)
Begin
    v1_r <= v1;
    v2_r <= v2;
    m_r <= v1_r*v2_r;
end
```

two-stage

```
always @(posedge clk)
Begin
    v1_r <= v1;
    v2_r <= v2;
    m0_r <= v1_r * v2_r;
    m_r <= m0_r
end
```

three-stage

Memory Stage

- ❑ The memory stage performs the following tasks:
 - Prepare 8-bit or 16-bit data for 32-bit buses
 - Setup the byte-selection signals
 - Detect memory alignment exceptions
 - Write a word to the data memory
- ❑ Aquila only supports aligned accesses for 16- or 32-bit data since non-aligned access takes multiple cycles
 - Generate misaligned exception if necessary
- ❑ Passing prior signals to the writeback stage
 - Including PC, instruction validity & exception signals
 - These signals are registered at the input of the writeback stage, instead of the output of the memory stage

Writeback Stage

- ❑ The Writeback unit performs physical write back to registers
 - The source may come from the Execute unit (through the Memory unit) or the data memory
 - The destination may be the register file, the CSR file, or the Forwarding unit

- ❑ Functions:
 - Data write back to registers
 - Perform sign-extension if necessary
 - Handles exception

Atomic Unit

- ❑ Atomic unit sits between the data cache and the DDRx memory controller, intercepting all memory requests
 - Each hardware thread (HART) has an ID
- ❑ Functions:
 - If the memory request is not from atomic instructions, bypass
 - If it is a lock-based AMO request, execute the operations
 - If it is a lock-free LR/SC request, manage the reservation table (one entry per HART):
 - Register the reservation for LR by its hart ID
 - Check the reservation and allow the first SC to cancel all other ID's reservations

Application Processors

- ❑ An application processor is an SoC contains:
 - Processor cores (ISA-based CPU/GPU/DSP)
 - Application-specific accelerators (none-ISA based)
 - audio/video codecs
 - image processing cores
 - machine-learning/AI cores
 - other application-specific cores
 - On-chip memories
 - Clock-domain infrastructure components
- ❑ ISA does not define an application processor

User-Defined Instructions

- ❑ An application processor based on the RISC-V ISA may define its own application-specific instructions
 - For 32-bit ISA, two of the 7-bit opcode are reserved for users: 0001011 (custom-0) and 0101011 (custom-1)
 - The rest of the 25-bit pattern can be defined arbitrarily
 - It is recommended to stick to the R/I/S formats

	31	25	24	20	19	15	14	12	11	7	6	0
R	funct7				rs2	rs1	func3	rd	opcode			
I	imm[11:0]					rs1	func3	rd	opcode			
S	imm[11:5]				rs2	rs1	func3	imm[4:0]	opcode			

- Programming of user-defined instructions can be done using inline assembly

Integration Issues of Multiple Cores

- ❑ Different cores often runs at different clock rates, and must exchange data and control information
- ❑ Hardware synchronization mechanisms
 - Atomic operations for heterogeneous cores
 - Data sharing
 - Unified coherent cache
 - Multi-port memory
- ❑ Clock-Domain Crossing (CDC) schemes
 - Shift registers
 - Asynchronous FIFOs

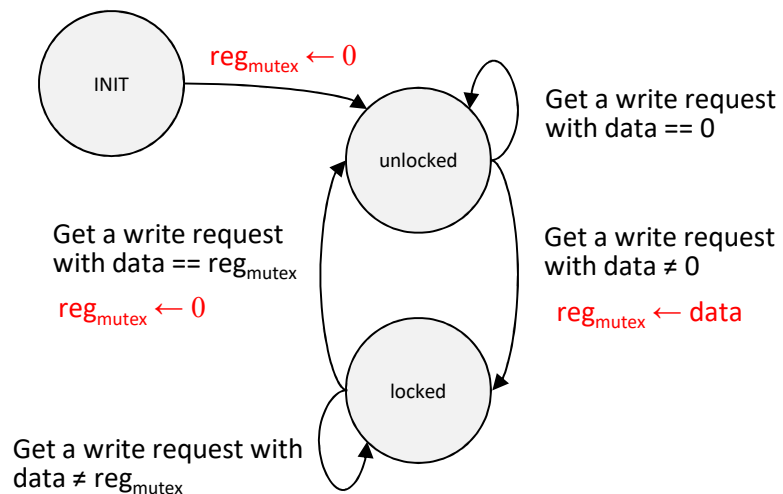
Atomicity cross Heterogeneous Cores

- ❑ The atomic instructions defined in ISA does not work across heterogeneous cores
- ❑ To guarantee atomic test-and-set, we can either use a coherent bus protocol, or hardware mutexes
- ❑ Atomic bus protocols can support
 - Bus locking
 - Multi-cycle read-and-write transactions (RMW cycle)


Hardware Mutex

- ❑ A HW mutex is a special memory device that “conditionally” accepts write requests
 - Suitable for synchronization between heterogeneous cores
 - Drawback: limited number of mutexes

- ❑ Hardware mutex controller:




- ❑ System diagram:



Cache Coherence Interconnect

- ❑ If each processor has its own data cache, there would be data coherence issues when two μ Ps share data:

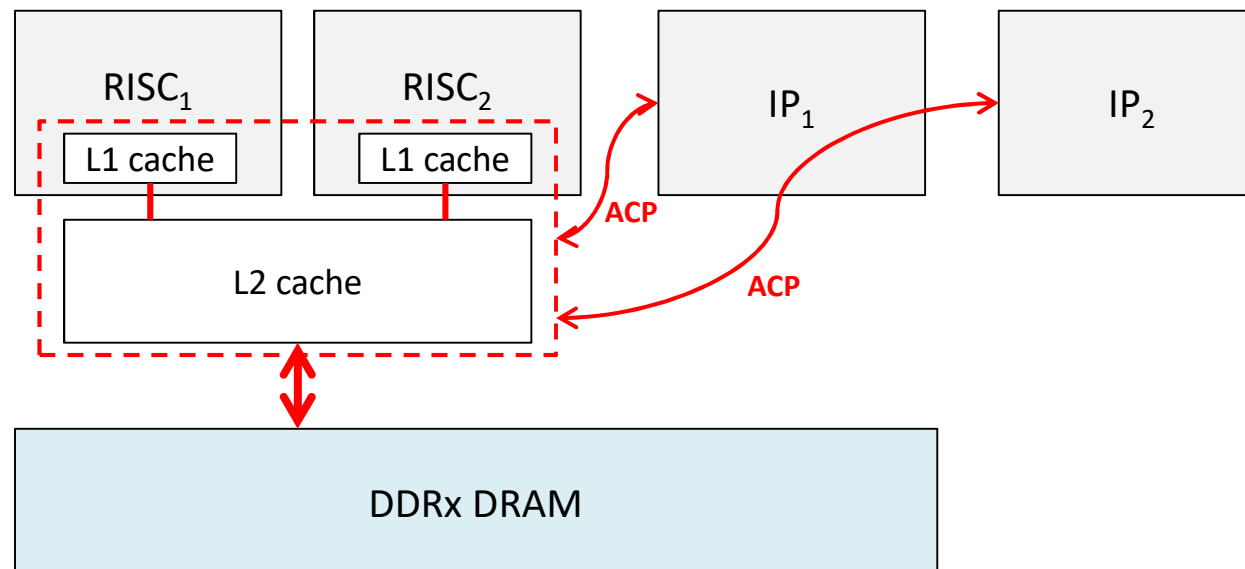


Cache Coherence Protocols

- ❑ If only a centralized cache controller is used, there is no cache coherence issue
 - When multiple cache controllers are used, there are two ways to guarantee cache coherence: snooping- and directory-based
- ❑ Snooping-based
 - All cache controllers of all processors monitor the bus traffics
 - A controller updates/invalidates its cache lines if they are modified by other controllers
- ❑ Directory-based
 - The shared cache line's state flags is in a common directory
 - A cache controller, when accessing a shared cache line must check the directory to ensure cache coherency

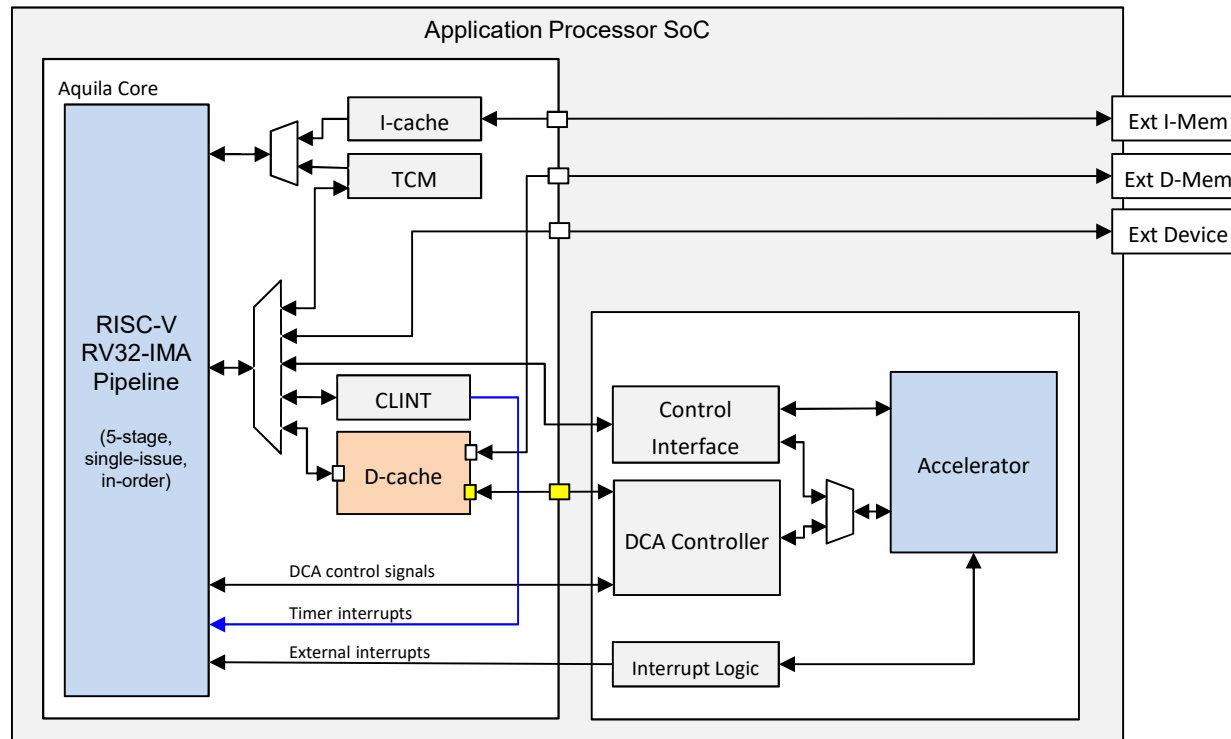
Heterogeneous Sharing of Caches

- ❑ For heterogeneous cores to share a cache:
 - General-purposes coherent bus protocols (e.g. AXI ACP)
 - Custom-made Direct-Cache Access controllers (DCA)
- ❑ Example: ARM AXI bus has the Accelerator Coherence Port (ACP) for accelerators to share the CPU cache



Direct Cache Access (DCA)

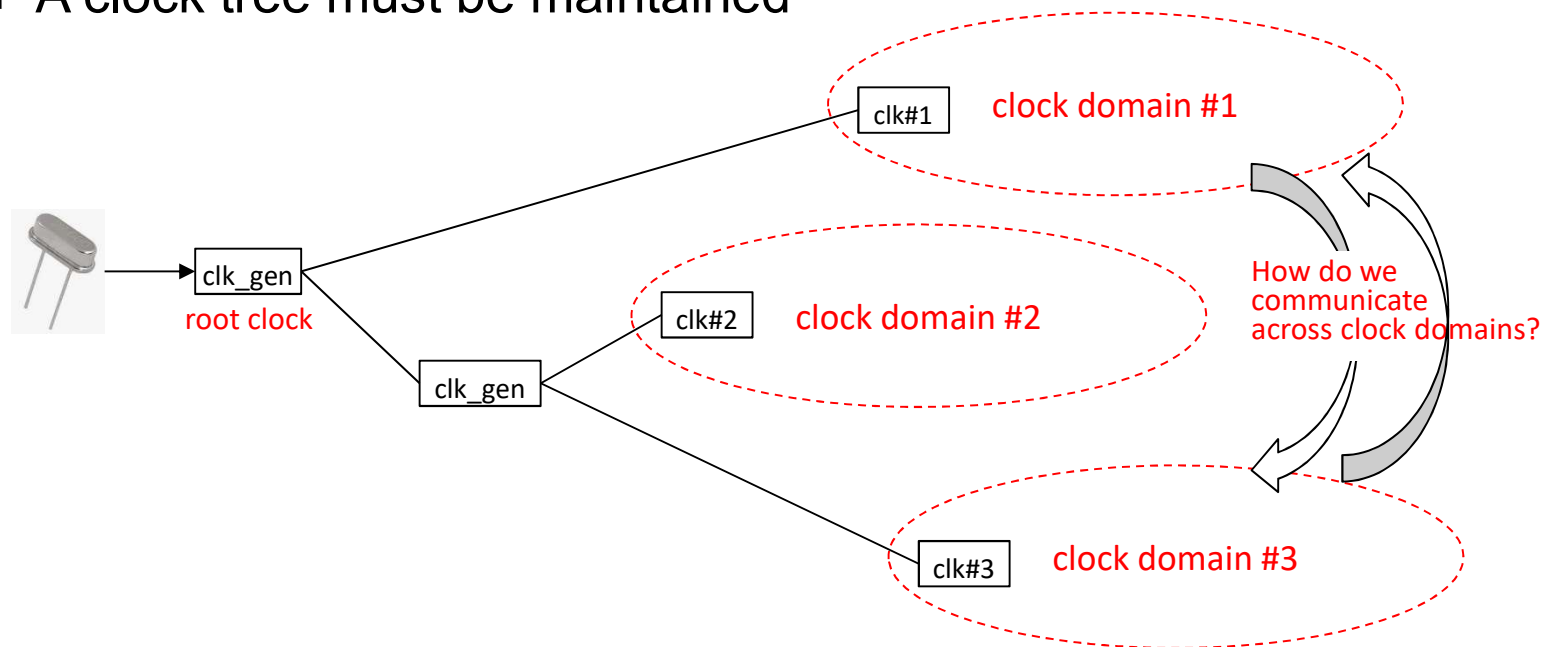
- ❑ The DCA concept was proposed by Intel[†]
 - Simulations show that transferring incoming packets directly into the CPU cache reduces misses by 11%



[†] R. Huggahalli, R. Iyer, S. Tetrack, "Direct Cache Access for High Bandwidth Network I/O," *Proc. of 32nd Int. Symp. on Computer Architecture (ISCA'05)*, Jun. 4-8, 2005.

Clock Domain Crossing

- ❑ A complex SoC usually covers several clock domains
 - A clock tree must be maintained



- The parameters of each clock must be set properly in the constraint files for the EDA tools to do static timing analysis

Clock Generation

- ❑ For FPGAs, there are several ways to generate a clocks based on an oscillator/clock inputs
 - By user logics – not recommended
 - Digital Clock Managers (DCM) – old DLL-style clock synthesizer
 - Phase-Lock Loops (PLL) – precise but less flexible synthesizer
 - Multi-Mode Clock Manager (MMCM) – combines the advantages of DCM and PLL
- ❑ Verilog code to invoke MMCM:

```
clk_wiz_0 Clock_Generator(  
    .clk_in1(sysclk_i), // Input clock from the oscillator  
    .clk_out1(clk),      // System clock for the Aquila SoC  
    .clk_out2(clk_166M), // Clock input to the MIG Memory controller  
    .clk_out3(clk_200M)  // DRAM Reference clock for MIG  
);
```

Clock Constraints

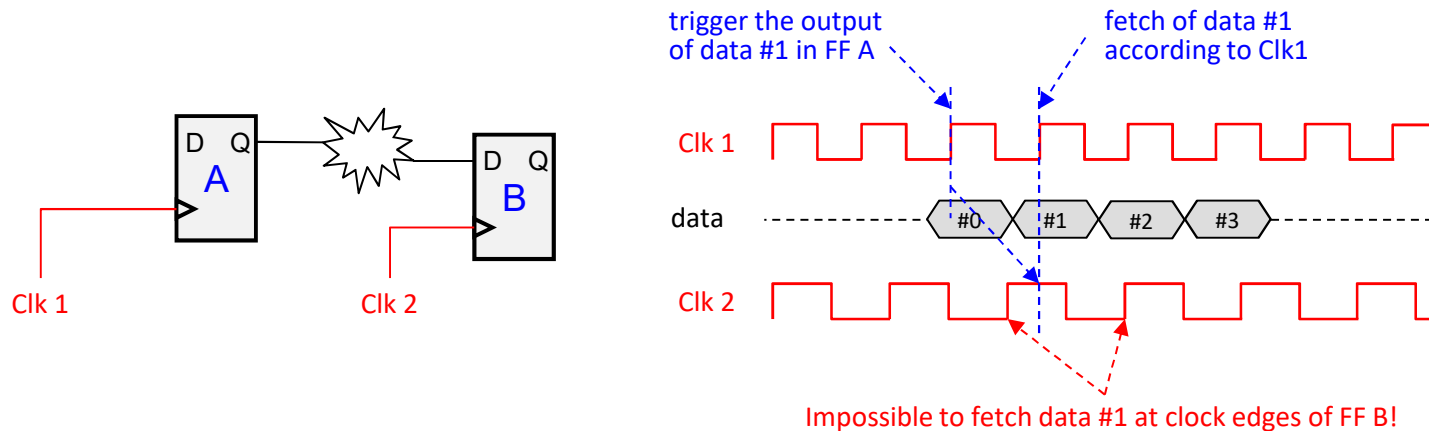
- ❑ If you use MMCM to generate clocks, the constraints for each clock will be automatically generated
 - **Don't** set constraints yourself, or you get design rule violation
- ❑ If you use user-logic to generate clocks, then you must specify the clock parameters in a constraint file:
 - For the clocks from the oscillator (i.e. the root clock):
 - For the derived clocks from user logic (clock multiplier/divider)

```
create_clock -name sysclk_i -period 10.000 -waveform {0.000 5.000} [get_ports sysclk_i];
```

```
create_generated_clock -period 24.000 -waveform {0.000 12.000} -name clk -source [get_pins usr_clock/clk_o]
```

CDC Issues


- ❑ When two flip-flops are driven by different clocks, data exchange can cause problems (a.k.a. meta-stability)



- ❑ Some simple solutions can be found in the white paper:
<https://zipcpu.com/blog/2017/10/20/cdc.html>


Passing Signals across Domains

- ❑ For general signal crossing, use asynchronous FIFOs:
 - BRAM on FPGA has asynchronous dual memory ports
 - The design of head/tail pointers is not trivial[†]



- Xilinx has a “fifo_generator” that creates asynchronous FIFOs

- ❑ For 1-bit signals, shift registers can be used:



[†] C. E. Cummings and P. Alfke, “Simulation and Synthesis Techniques for Asynchronous FIFO design,” Synopsys Users Group Meetings, San Jose, CA, 2002.

Discussions

- ❑ Application Processor design is the key differentiator for smartphone processors
 - Many processor are based on exactly the same ARM core, but they are different processors (e.g., from MediaTek, Qualcomm, or Samsung)
- ❑ For best performance, some accelerators must be tightly integrated with the μ arch of the processor core
 - Knowing the guts of the μ arch is important to design a high performing application processor