

HW#4 RTOS Analysis



Chun-Jen Tsai
NYCU
11/17/2023

Homework Goal

- ❑ In this homework, you will analyze the performance of a real-time OS (RTOS), FreeRTOS, for multithreading
- ❑ Your tasks:
 - Trace the OS kernel code and analyze how thread management and synchronization are done
 - Measure the context-switching overhead of the application (need to add counters in Aquila)
 - Measure the synchronization overhead of the application (better to add counters in Aquila)
- ❑ You should upload your report to E3 by 12/8, 17:00.
 - Report is 4 pages max, PDF format only. No demo this time.

FreeRTOS

- ❑ FreeRTOS is a C-based real-time operating system kernel for embedded devices
 - Developed by Richard Barry in 2003
 - Barry joined Amazon Web Services (AWS) and passed the stewardship of the project to AWS in 2017
 - The project adopts MIT License

- ❑ We used FreeRTOS v202111.00 in this HW:
 - All RISC-V unrelated sources are removed (way too big)
 - The original source available at <https://www.freertos.org/>
 - No need to modify source code of FreeRTOS for Aquila

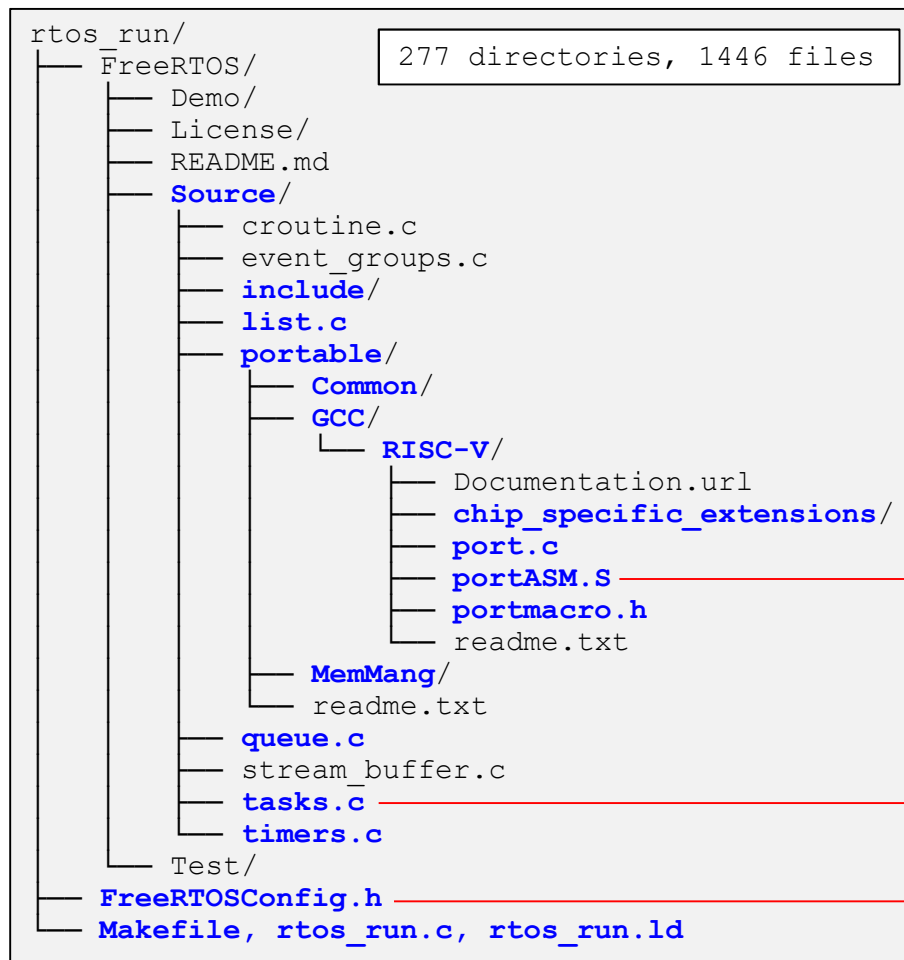
FreeRTOS Multithreading API

- ❑ Unlike other programming languages, C does not have a standard API for multithreading
 - ISO C11 has a multithreading API since 2011, but the most popular API is still the non-standard pthread API
- ❑ FreeRTOS Multithreading API is quite simple:

```
/* Two threads creation for FreeRTOS */  
  
int main(void)  
{  
    int prm1 = 1, prm2 = 2;  
    xTaskCreate(Task_Handler, "Task1", 256, (void *) &prm1, 3, NULL);  
    xTaskCreate(Task_Handler, "Task2", 256, (void *) &prm2, 4, NULL);  
    vTaskStartScheduler();  
}  
  
void Task_Handler(void *pvParam)  
{  
    for (int idx = 0; idx < 10/(int) *pvParam; idx++) {  
        printf("\nThis is Task%d.\n", (int) *pvParam);  
        vTaskDelay(1000/portTICK_PERIOD_MS); // sleep a while.  
    }  
  
    vTaskDelete(NULL); /* Thread ends, delete it from the task queue. */  
}
```

Target Application: rtos_run

❑ Download rtos_run.tgz from E3:



After "make", a build/ directory would contain the rtos_run.map and rtos_run.objdump files.

The executable rtos_run.elf will be in the top directory

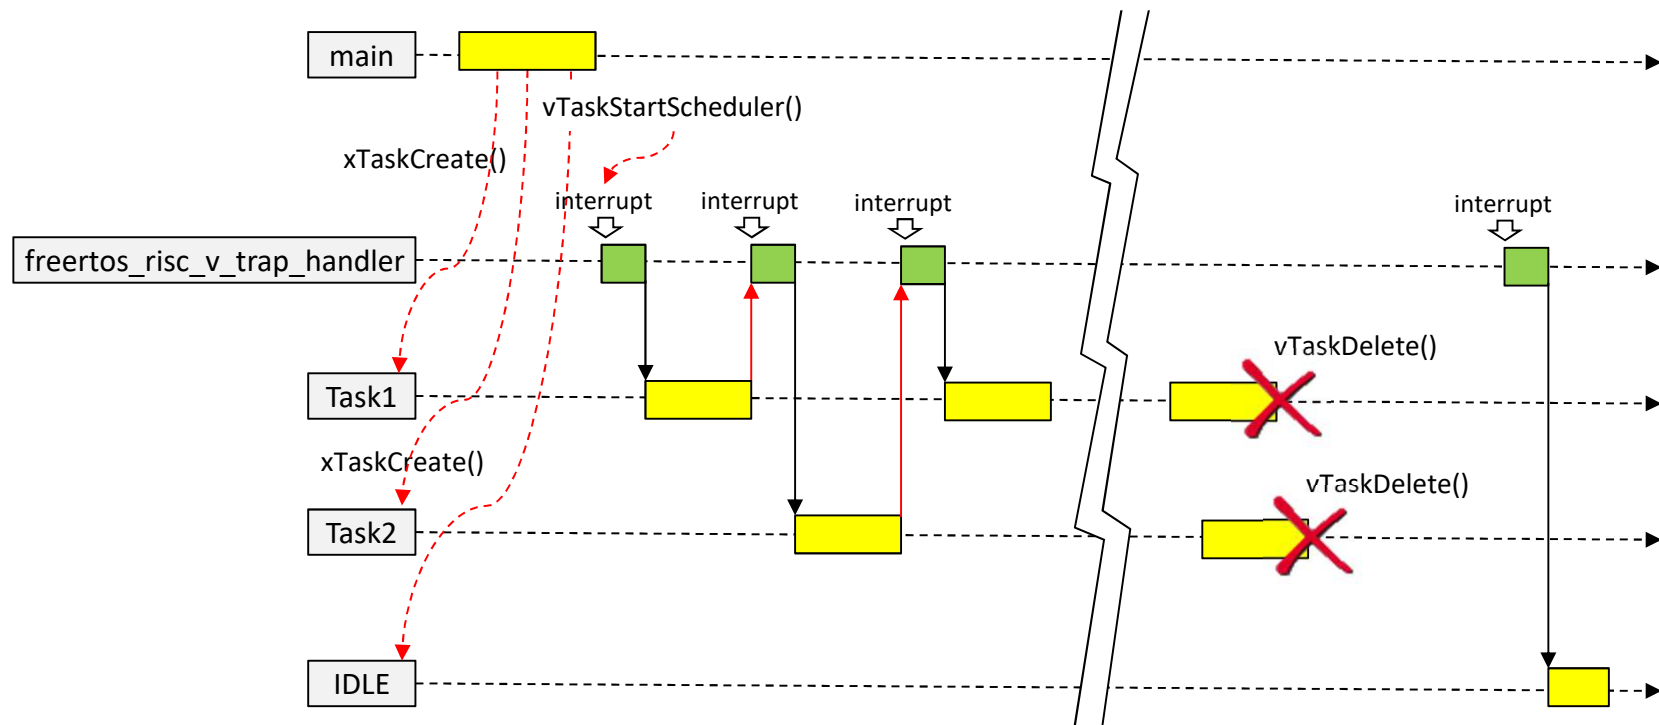
Timer interrupt handling routine.

Multi-thread managers.

FreeRTOS configurations for application.

Application Behavior of `rtos_run`

- There are three visible threads and two invisible threads in the `rtos_run` application:



Protection of Shared Resources

- ❑ In a preemptive multi-tasking OS, shared resources cannot be modified without protection:

```
volatile int
shared_counter = 0,
done = 0;

Task1()
{
    while (!done)
    {
        shared_counter++;
    }
}

Task2()
{
    while (!done)
    {
        shared_counter++;
    }
}
```

```
80000030 <Task1>:
80000030: lui    a3, 0x80008
80000034: lw     a5, 8(a3)    # <done>
80000038: bnez   a5, 80000054 <Task1+0x24>
8000003c: lui    a4, 0x80008
80000040: lw     a5, 12(a4)   # <shared_counter>
80000044: addi   a5, a5, 1
80000048: sw     a5, 12(a4)
8000004c: lw     a5, 8(a3)
80000050: beqz   a5, 80000040 <Task1+0x10>
80000054: ret
```

Default Execution

- ❑ When you compile and run the application you should see the following output
 - The result is bad because we did not enable mutex protection

```
=====
Copyright (c) 2019-2023, EISL@NYCU, Hsinchu, Taiwan.
The Aquila SoC is ready.
Waiting for an ELF file to be sent from the UART ...

Program entry point at 0x8000144C, size = 0xAE60.
-----

Task 1 start running ...

Task 2 start running ...


At the end, the shared counter = 10000
Task1 local counter = 10000
Task2 local counter = 9263
Task1 counter + Task2 counter != Shared counter, the counter is corrupted.
```


Application with Mutex Protection

- ❑ With mutex protection of the shared variable, the output would be good
 - Made the change: “#define USE_MUTEX 1” in `rtos_run.c`

```
=====
Copyright (c) 2019-2023, EISL@NYCU, Hsinchu, Taiwan.
The Aquila SoC is ready.
Waiting for an ELF file to be sent from the UART ...

Program entry point at 0x800014C4, size = 0xAE60.
-----

Task 1 start running ...

Task 2 start running ...


At the end, the shared counter = 10000
Task1 local counter = 3065
Task2 local counter = 6935
The shared counter is protected well.
```

Context-Switching Overhead

- ❑ A preemptive multi-tasking OS uses timer interrupts to assign CPU usage from one thread to the other
 - The context-switching overhead is inversely proportional to the time quantum (slice)
 - The default of FreeRTOS time quantum is 10 msec
- ❑ To measure the context-switching overhead, you must count the number of cycles between:
 - A timer interrupt arrives
 - A new thread begins execution
- ❑ You can change the time quantum size down to 5 msec to see its impact on the overhead

The Timer Interrupt Device

- ❑ In Aquila, the module `clint` is used to provide timer interrupts and software interrupts
- ❑ `Clint` has three registers
 - `mtime`: 64-bit counter of timer ticks
 - `mtimecmp`: Upper-threshold to trigger a timer interrupt
 - `msip`: a 32-bit register to trigger a software interrupt
- ❑ FreeRTOS will update `mtimecmp` to setup the next context-switch time (based on time quantum duration)

Changing Time Quantum

- ❑ In FreeRTOS, time quantum is configured by the header file: `FreeRTOSConfig.h`
 - The default time quantum is 10 msec:

```
...  
  
#define CLINT_CTRL_ADDR          ( 0xF0000000UL )  
#define configMTIME_BASE_ADDRESS ( CLINT_CTRL_ADDR + 0x0UL )  
#define configMTIMECMP_BASE_ADDRESS ( CLINT_CTRL_ADDR + 0x8UL )  
  
#define configUSE_PREEMPTION      1  
#define configUSE_IDLE_HOOK      0  
#define configUSE_TICK_HOOK      1  
#define configCPU_CLOCK_HZ       ( ( uint32_t ) ( 41666667 ) )  
#define configTICK_RATE_HZ       ( ( TickType_t ) 100 )  
#define configMAX_PRIORITIES     ( 7 )  
#define configMINIMAL_STACK_SIZE ( ( uint32_t ) 100 )  
#define configTOTAL_HEAP_SIZE    ( ( size_t ) ( 12 * 1024 ) )  
  
...
```

Mutex for Synchronization

- ❑ A mutex is a variable to indicate two states: “locked” and “unlocked” of a shared resource:

```
int mutex;
```

```
mutex_take(mutex);
```

code that uses the shared resource.

```
mutex_give(mutex);
```

Execution blocked here if the mutex has been taken by other threads.

- ❑ There are several approaches to implement a mutex: ISA-independent SW, ISA-dependent SW, or hardware

Mutex Implementation

- ❑ Software mutex implementation techniques
 - Software algorithms (e.g. the Peterson's algorithm)
 - Drawback: time-consuming
 - Atomic test-and-set
 - Drawback: less efficient for multi-core systems
 - Conditional load-store
 - Drawback: only supported by new ISAs & CPUs
- ❑ Hardware mutex approach
 - A HW mutex is a device that contains a list of mutex registers
 - An unlocked register has zero in it
 - Each thread write their ID to the register to lock the mutex
 - Each register conditionally accepts the write requests
 - Suitable for synchronization even between HW and SW

Peterson's Mutex Algorithm

- ❑ Peterson's algorithm[†] guarantees exclusive accesses to a shared resource among n threads (running on n cores) without special assembly instructions
- ❑ A two-thread version is as follows:

CPU 0

```
/* trying protocol for T1 */  
Q1 = true; /* request to enter */  
TURN = 1; /* who's turn to wait */  
wait until not Q2 or TURN == 2;  
Critical Section;  
/* exit protocol for T1 */  
Q1 = false;
```

CPU 1

```
/* trying protocol for T2 */  
Q2 = true; /* request to enter */  
TURN = 2; /* who's turn to wait */  
wait until not Q1 or TURN == 1;  
Critical Section;  
/* exit protocol for T2 */  
Q2 = false;
```

[†] G. L. Peterson, "Myth about the Mutual Exclusion Problem," *Information Processing Letters*, **12**, no 3, June 30, 1981.

Test-and-Set Atomic Instructions

- ❑ For synchronization, a thread must execute the following code before entering a critical section:

```
int mutex; /* '0' means unlocked, '1' means locked */  
while ( test_and_set(mutex) == 1) /* busy waiting */  
    Code that uses the shared resource.  
mutex = 0;
```

- ❑ A 'SWAP' instruction (`amoswap.w` in RSIC-V) can be used to implement the test-and-set function:

```
int test_and_set(int mutex)  
{  
    temp = mutex;  
    mutex = 1;  
    return temp;  
}
```

The first two lines cannot be interrupted during execution!

Conditional Load/Store Instructions

- ❑ Conditional load/store allows atomic operation without locking the buses
 - In RISC-V, we have LR/SC instructions
 - In ARM, we have LDREX/STREX instructions
- ❑ In this HW, you should try to use the atomic instructions (either lock-based or lock-free) to implement mutex and measure the overhead
 - The atomic operations are implemented in `atomic_unit.v`

Example Code:

❑ Mutex take:

```
asm volatile ("lui t0, %hi(lock_addr)");  
asm volatile ("lw t3, %lo(lock_addr)(t0)");  
asm volatile ("li t0, 1");  
asm volatile ("0:");  
asm volatile ("lw t1, (t3)");  
asm volatile ("bnez t1, 0b");  
asm volatile ("amoswap.w.aq t1, t0, (t3)");  
asm volatile ("bnez t1, 0b");
```

❑ Mutex give:

```
asm volatile ("lui t0, %hi(lock_addr)");  
asm volatile ("lw t3, %lo(lock_addr)(t0)");  
asm volatile ("amoswap.w.rl x0, x0, (t3)");
```

Mutex Take/Give in FreeRTOS

- ❑ Note that, for FreeRTOS, most of the overhead in mutex take and give operations are spent in handling priority inversion
 - Queue structures are used to avoid a low priority task to block high priority tasks
 - We do not have different priorities for the two threads in the sample program
- ❑ Removing priority inversion processing reduces the synchronization overhead significantly, but it is not the right thing to do.

Synchronization Overhead

- ❑ In the sample application, we use two synchronization schemes of FreeRTOS: critical sections and mutex
 - Critical sections are used to protect the UART device so that different threads can print concurrent messages properly
 - A mutex is use to protect a shared variable so that both threads can modify the variable without corrupting it
 - In FreeRTOS, mutex is a special type of semaphore.
- ❑ You should measure the overhead of these synchronization schemes and put that in your report
 - The cycles to enter & leave a critical section
 - The cycles to take & give a mutex

Comments on the Homework

- ❑ The key point of this homework is RTOS multi-threading analysis
 - Performance optimization is not required
- ❑ Your grade will be totally based on your analysis of the RTOS behavior:
 - The context switching behavior analysis
 - The algorithmic description
 - Context switching overhead vs. time quantum
 - The synchronization behavior analysis
 - The algorithmic description
 - The overhead (cycles required for mutex task & give, respectively)
 - Any additional analyses you can think of regarding multithreading (e.g. impact on I\$/D\$)